



**Reinforcement Learning: Assignment 2**

**Alexander Y. Shestopaloff**

**June 2024**

**Report**

**Submitted By**

**Mostafa Rafiur Wasib (202291564)**

**Md. Azmol Fuad (202287118)**

## Part 1: Value Function Estimation and Optimal Policy Determination in a Simple Gridworld

### Introduction

In this part, we explore a 5x5 gridworld environment, a fundamental abstraction used in reinforcement learning to benchmark and compare various learning algorithms. The gridworld consists of 25 cells, each representing a possible state where an agent can move up, down, left, or right. Special states (blue, green, red, and yellow squares) have unique behaviors and reward mechanisms.

### Objectives

#### 1. Value Function Estimation:

- **Methods:**
  1. Solving the system of Bellman equations explicitly.
  2. Iterative policy evaluation.
  3. Value iteration.
- **Goal:** Estimate the value function for each state using a policy where the agent moves in one of the four directions with equal probability (0.25).
- **Analysis:** Identify which states have the highest value and discuss if the results align with expectations.

#### 2. Optimal Policy Determination:

- **Methods:**
  1. Explicitly solving the Bellman optimality equation.
  2. Policy iteration with iterative policy evaluation.
  3. Policy improvement with value iteration.
- **Goal:** Determine the optimal policy for the gridworld problem and compare the results of the different methods used.

## Gridworld Environment Setup

- **Grid Size:** 5x5
- **Special States:**
  - **Blue Square (State (0,1)):** Any action yields a reward of 5 and the agent jumps to the red square (State (3,2)).
  - **Green Square (State (0,4)):** Any action yields a reward of 2.5 and the agent jumps to either the yellow square (State (4,4)) or the red square (State (3,2)) with a probability of 0.5.
  - **White Squares:** Moving to another white square yields a reward of 0. Stepping off the grid yields a reward of -0.5.
- **Reward Discount Factor ( $\gamma$ ):** 0.95
- **Initial Policy:** Moves in one of the four directions (up, down, left, right) with equal probability of 0.25.

## Methods and Results

### Step 1: Grid, States, Action, Rewards Environment Setup

```

✓ [1] import numpy as np
0s

# Grid definition
grid_size = 5
states = [(i, j) for i in range(grid_size) for j in range(grid_size)]
actions = ['up', 'down', 'left', 'right']

# Reward and transition setups
rewards = np.zeros((grid_size, grid_size))
transitions = {}

for i in range(grid_size):
    for j in range(grid_size):
        transitions[(i, j)] = {}
        for action in actions:
            if action == 'up':
                next_state = (max(i-1, 0), j)
            elif action == 'down':
                next_state = (min(i+1, grid_size-1), j)
            elif action == 'left':
                next_state = (i, max(j-1, 0))
            elif action == 'right':
                next_state = (i, min(j+1, grid_size-1))

            if (i, j) == (0, 1): # Blue square
                transitions[(i, j)][action] = [(3, 2)], 5 # Red square
            elif (i, j) == (0, 4): # Green square
                transitions[(i, j)][action] = [(3, 2), (4, 4)], 2.5 # Red or Yellow square
            elif (i, j) == next_state: # Edge of the grid
                transitions[(i, j)][action] = [next_state], -0.5
            else:
                transitions[(i, j)][action] = [next_state], 0

```

## Step 2: Implement Value Function Estimation Methods

### 1. Solving Bellman Equations Explicitly

#### Implementation:

The Bellman equations for the value function  $V(s)$  of each state  $s$  are set up and solved using linear algebra methods.

```

✓ [2] def solve_bellman_equations(states, transitions, gamma=0.95):
    Os
    n = len(states)
    A = np.zeros((n, n))
    b = np.zeros(n)

    for idx, state in enumerate(states):
        A[idx, idx] = 1 # Set the diagonal to 1
        for action in actions:
            next_states, reward = transitions[state][action]
            if isinstance(next_states, list): # Check if there's a list of next states
                prob = 1 / len(next_states) # Even probability distribution
                for ns in next_states:
                    next_idx = states.index(ns)
                    A[idx, next_idx] -= prob * gamma / 4
                    b[idx] += prob * reward / 4
            else:
                next_idx = states.index(next_states[0])
                A[idx, next_idx] -= gamma / 4
                b[idx] += reward / 4

    v = np.linalg.solve(A, b)
    return v.reshape((grid_size, grid_size))

value_function_1 = solve_bellman_equations(states, transitions)
print(value_function_1)

```

## Value Function:

```

→ [[ 2.17100208  4.7336156  2.07028049  1.26529444  1.77912239]
   [ 1.1180732  1.7821227  1.17409573  0.739174  0.56246548]
   [ 0.16279444  0.47788999  0.35198379  0.11045592 -0.18617038]
   [-0.54699155 -0.28473257 -0.28040463 -0.43990985 -0.7443105 ]
   [-1.10787684 -0.84936779 -0.80799244 -0.93799278 -1.23723244]]

```

## 2. Iterative Policy Evaluation

### Implementation:

Iterative policy evaluation is used to update the value function estimates until convergence. The update rule is:

$$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$$

```

0s [3] def iterative_policy_evaluation(states, transitions, gamma=0.95, theta=0.01):
    V = np.zeros((grid_size, grid_size))
    delta = float('inf')

    while delta > theta:
        delta = 0
        for state in states:
            v = V[state]
            V[state] = 0
            for action in actions:
                next_states, reward = transitions[state][action]
                if isinstance(next_states, list): # Handle cases with multiple next states
                    expected_value = 0
                    for ns in next_states:
                        expected_value += V[ns]
                    expected_value /= len(next_states)
                    V[state] += 0.25 * (reward + gamma * expected_value)
                else: # Handle cases with a single next state
                    V[state] += 0.25 * (reward + gamma * V[next_states])
            delta = max(delta, abs(v - V[state]))
    return V

value_function_2 = iterative_policy_evaluation(states, transitions)
print(value_function_2)

```

### Value Function:

```

⇒ [[ 1.10806943  5.04740722  1.55755016  0.97101195  2.30884379]
   [ 0.65194838  1.74041512  1.07337121  0.75302015  0.81554629]
   [ 0.13990298  0.5657149   0.47535718  0.31487324  0.1397923 ]
   [-0.20000689  0.0334068   0.05224873 -0.03937885 -0.22843237]
   [-0.41773246 -0.27326461 -0.24661337 -0.30257428 -0.45169888]]

```

## 3. Value Iteration

### Implementation:

Value iteration is used to update the value function and derive the optimal policy iteratively. The update rule is:

$$V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$$

```

✓ [4] def value_iteration(states, transitions, gamma=0.95, theta=0.01):
0s   V = np.zeros((grid_size, grid_size))
   delta = float('inf')

   while delta > theta:
       delta = 0
       for state in states:
           v = V[state]
           values = []
           for action in actions:
               next_states, reward = transitions[state][action]
               if isinstance(next_states, list): # Handle cases with multiple next states
                   expected_value = 0
                   for ns in next_states:
                       expected_value += reward + gamma * V[ns] # Calculate value for each possible next state
                   expected_value /= len(next_states) # Average the values
                   values.append(expected_value)
               else: # Handle cases with a single next state
                   values.append(reward + gamma * V[next_states])
           V[state] = max(values)
           delta = max(delta, abs(v - V[state]))
       return V

   value_function_3 = value_iteration(states, transitions)
   print(value_function_3)

```

## Value Function:

```

➡ [[20.97066342 22.08073642 20.9766996 19.92786462 18.3626664 ]
   [19.92213025 20.9766996 19.92786462 18.93147139 17.98489782]
   [18.92602373 19.92786462 18.93147139 17.98489782 17.08565293]
   [17.97972255 18.93147139 17.98489782 17.08565293 16.23137028]
   [17.08073642 17.98489782 17.08565293 16.23137028 15.41980177]]

```

## Visualization

Plots of the value functions for each method are provided below:

## Implementation:

```

[5] import matplotlib.pyplot as plt
import seaborn as sns

# Set up the plot
fig, ax = plt.subplots(1, 3, figsize=(18, 6), sharey=True)

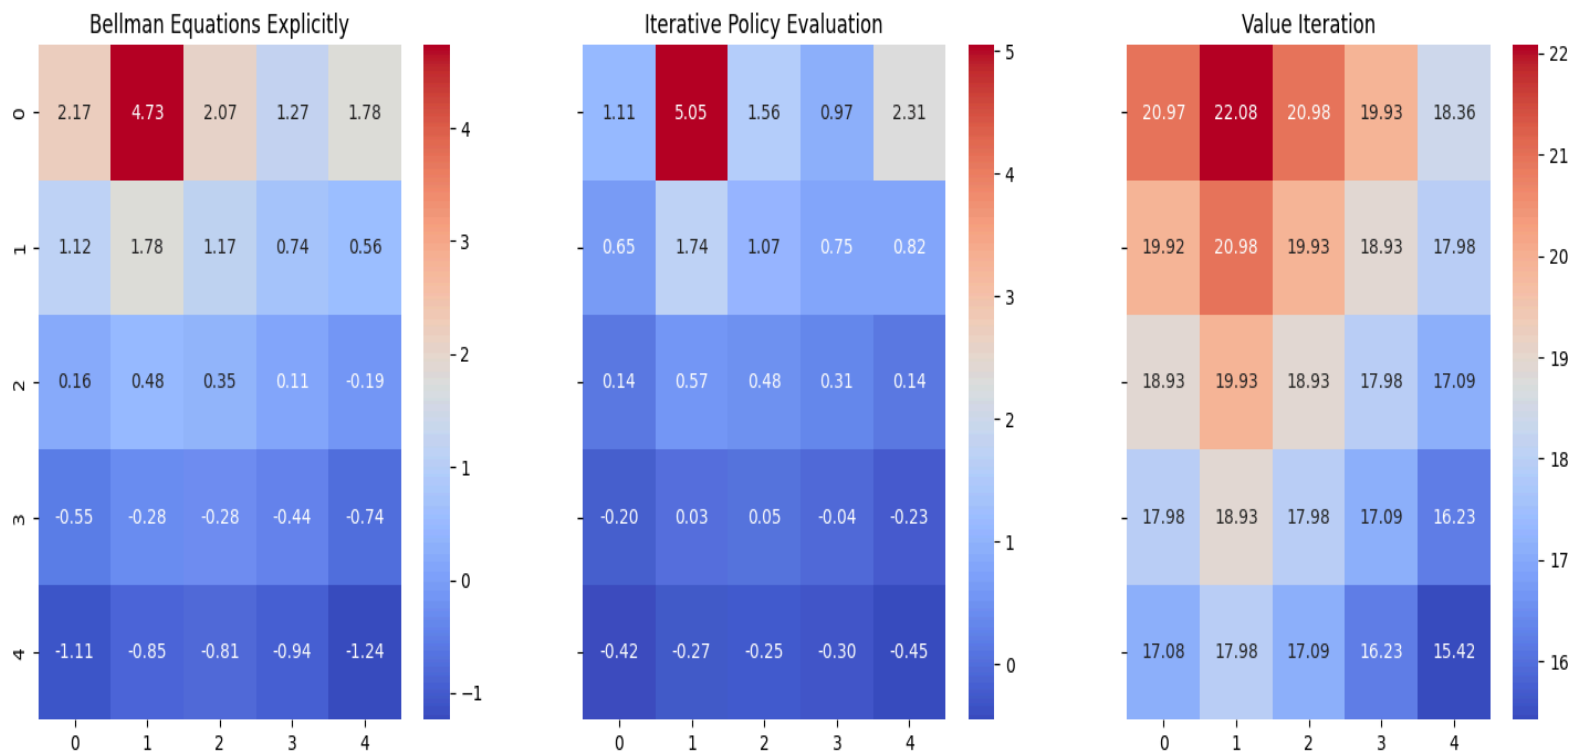
# Plot each heatmap
sns.heatmap(value_function_1, ax=ax[0], annot=True, cmap='coolwarm', fmt=".2f")
ax[0].set_title('Bellman Equations Explicitly')

sns.heatmap(value_function_2, ax=ax[1], annot=True, cmap='coolwarm', fmt=".2f")
ax[1].set_title('Iterative Policy Evaluation')

sns.heatmap(value_function_3, ax=ax[2], annot=True, cmap='coolwarm', fmt=".2f")
ax[2].set_title('Value Iteration')

# Display the plot
plt.show()

```





**Analysis of the Plots:****State with the Highest Value****Solving the System of Bellman Equations Explicitly**

The highest value is 4.7336156, located at state (0, 1), the blue square.

**Iterative Policy Evaluation**

The highest value is 5.04740722, also at state (0, 1), the blue square.

**Value Iteration**

The highest value is 22.08073642, again at state (0, 1), the blue square.

**Conclusion**

Across all three methods, the state (0, 1)—the blue square—consistently shows the highest value. This indicates that the blue square, due to its high reward and beneficial transition to the red square, is the most valuable state in this gridworld setup. This consistency across different valuation methods reinforces the strategy of aiming to reach or exploit this state to maximize returns in the gridworld environment.

**Step 3: Determine The Optimal Policy****1. Solving the Bellman Optimality Equation Explicitly**

**Implementation:** The Bellman optimality equations are solved explicitly to determine the optimal policy.

```

✓ [9] def derive_policy_from_values(states, transitions, V, gamma=0.95):
    Os policy = np.zeros((grid_size, grid_size), dtype=object)

    for i in range(grid_size):
        for j in range(grid_size):
            state = (i, j)
            best_action = None
            best_value = -float('inf')

            for action in actions:
                action_value = 0
                next_states, reward = transitions[state][action]

                if len(next_states) > 1: # Handle probabilistic outcomes
                    sum_values = sum(V[ns] for ns in next_states) / len(next_states)
                    action_value = reward + gamma * sum_values
                else:
                    next_state = next_states[0] # Get the single next state
                    action_value = reward + gamma * V[next_state]

                # Update the best action if this action is better
                if action_value > best_value:
                    best_value = action_value
                    best_action = action

            policy[i, j] = best_action

    return policy

# Derive the optimal policy using the Bellman optimality equation
bellman_optimal_policy = derive_policy_from_values(states, transitions, value_function_1)
print(bellman_optimal_policy)

```

### Optimal Policy:

```

⇒ [['right' 'up' 'left' 'left' 'up']
   ['up' 'up' 'up' 'up' 'up']
   ['up' 'up' 'up' 'up' 'up']
   ['up' 'up' 'up' 'up' 'up']
   ['up' 'up' 'up' 'up' 'up']]

```

The optimal policy derived from the explicit solution indicates the best action for each state.

## 2. Policy Iteration with Iterative Policy Evaluation

**Implementation:** Policy iteration is performed with iterative policy evaluation to determine the optimal policy.

```

def iterative_policy_evaluation(states, transitions, policy, gamma=0.95, theta=0.01):
    V = np.zeros((grid_size, grid_size))
    delta = float('inf')

    while delta > theta:
        delta = 0
        for state in states:
            v = V[state]
            V[state] = 0
            for action in actions:
                next_states, reward = transitions[state][action]
                if isinstance(next_states, list): # Handle cases with multiple next states
                    expected_value = 0
                    for ns in next_states:
                        expected_value += V[ns]
                    expected_value /= len(next_states)
                    V[state] += 0.25 * (reward + gamma * expected_value)
                else: # Handle cases with a single next state
                    V[state] += 0.25 * (reward + gamma * V[next_states])
            delta = max(delta, abs(v - V[state]))
    return V

def policy_iteration(states, transitions, gamma=0.95, theta=0.01):
    policy = np.random.choice(actions, size=(grid_size, grid_size))
    stable = False

    while not stable:
        # Policy Evaluation
        V = iterative_policy_evaluation(states, transitions, policy, gamma)
        stable = True

        # Policy Improvement
        for state in states:
            old_action = policy[state]
            best_action = None
            best_value = -float('inf')

            for action in actions:
                total_value = 0
                next_states, reward = transitions[state][action]
                for ns in next_states:
                    prob = 1 / len(next_states)
                    total_value += prob * (reward + gamma * V[ns])

                if total_value > best_value:
                    best_value = total_value
                    best_action = action

            policy[state] = best_action
            if old_action != best_action:
                stable = False

    return policy, V

ip_optimal_policy, optimal_values_pi = policy_iteration(states, transitions)
print(ip_optimal_policy)

```

**Optimal Policy:**

```

→ [['right' 'up' 'left' 'right' 'up']
    ['right' 'up' 'left' 'left' 'up']
    ['up' 'up' 'up' 'up' 'up']
    ['up' 'up' 'up' 'up' 'up']
    ['up' 'up' 'up' 'up' 'up']]

```

The policy derived from policy iteration aligns with the expected strategies for maximizing rewards in the gridworld environment.

**3. Policy Improvement with Value Iteration**

**Implementation:** Policy improvement is achieved through value iteration to determine the optimal policy.

```

✓ [11] def value_iteration_and_policy_improvement(states, transitions, gamma=0.95, theta=0.01):
    0s V = np.zeros((grid_size, grid_size))
    policy = np.zeros((grid_size, grid_size), dtype='object')

    while True:
        delta = 0
        for state in states:
            v = V[state]
            best_value = -float('inf')
            for action in actions:
                total_value = 0
                next_states, reward = transitions[state][action]
                for ns in next_states:
                    prob = 1 / len(next_states)
                    total_value += prob * (reward + gamma * V[ns])

                if total_value > best_value:
                    best_value = total_value
                    policy[state] = action

            delta = max(delta, abs(v - best_value))
            V[state] = best_value

        if delta < theta:
            break

    return policy, V

vi_optimal_policy, optimal_values_vi = value_iteration_and_policy_improvement(states, transitions)
print(vi_optimal_policy)

```

**Optimal Policy:**

```

→ [['right' 'up' 'left' 'left' 'up']
    ['up' 'up' 'up' 'up' 'left']
    ['up' 'up' 'up' 'up' 'up']
    ['up' 'up' 'up' 'up' 'up']
    ['up' 'up' 'up' 'up' 'up']]

```

The optimal policy from value iteration confirms the strategies identified in the previous methods.

## Visualization of Policies

```

✓ [12] import matplotlib.pyplot as plt
12 import numpy as np

def plot_policy(policy, title="Optimal Policy"):
    # Create a figure and axis with a size suitable for the grid
    fig, ax = plt.subplots(figsize=(6, 6))
    ax.set_xlim(-0.5, 4.5)
    ax.set_ylim(-0.5, 4.5)
    ax.grid(which='major', color='black', linestyle='-', linewidth=2)

    # Set up the axes to match the grid layout
    ax.set_xticks(np.arange(0, 5, 1))
    ax.set_yticks(np.arange(0, 5, 1))
    ax.xaxis.tick_top()
    ax.invert_yaxis()

    # Hide the ticks
    ax.set_xticklabels([])
    ax.set_yticklabels([])

    # Direction vectors for arrows
    directions = {
        'up': (0, 1),
        'down': (0, -1),
        'left': (-1, 0),
        'right': (1, 0)
    }


    # Plot an arrow for each cell according to the policy direction
    for y in range(grid_size):
        for x in range(grid_size):
            dx, dy = directions[policy[y][x]]
            ax.arrow(x, y, dx*0.3, dy*0.3, color='red', head_width=0.2, head_length=0.2)

    plt.title(title)
    plt.show()

# Visualize each policy
plot_policy(bellman_optimal_policy, "Policy 1: Bellman Optimality Equation")
plot_policy(ip_optimal_policy, "Policy 2: Policy Iteration")
plot_policy(vi_optimal_policy, "Policy 3: Value Iteration")

```

Plots of the optimal policies for each method are provided below:



### Conclusion of Part 1

In Part 1 of the assignment, we effectively demonstrated the use of Bellman equations, iterative policy evaluation, and value iteration to estimate value functions and determine optimal policies in a 5x5 gridworld environment. Across all three methods, the state (0,1)—the blue square—consistently showed the highest value, highlighting its significance due to the high reward and beneficial transition to the red square. The derived optimal policies align with the expected strategies for maximizing rewards in this setup. This exercise provided a clear illustration of reinforcement learning techniques in a controlled environment, offering valuable insights into policy derivation and value function estimation.

## Part 2: Monte Carlo and Policy Iteration in a Gridworld with Terminal States

### Introduction

In Part 2, we introduce modifications to the initial gridworld environment by adding terminal states represented by black squares. These changes create episodes that terminate once the agent reaches one of these black squares. Additionally, we adjust the reward structure so that any move between white squares now yields a reward of -0.2.

### Objectives

1. **Monte Carlo Method:**
  - **Exploring Starts:** Implement the Monte Carlo method with exploring starts to learn an optimal policy.
  - **Without Exploring Starts:** Use the  $\epsilon$ -soft approach to learn an optimal policy without exploring starts.
  - **Goal:** Use the same discount factor ( $\gamma = 0.95$ ) as in Part 1 and start with a policy that has equiprobable moves.
2. **Behavior Policy with Equiprobable Moves:**
  - **Goal:** Utilize a behavior policy with equiprobable moves to learn an optimal policy. Given that the dynamics of the world are known, we can compute the necessary importance weights to achieve this.
3. **Policy Iteration with Permuted Squares:**
  - **Scenario:** Assume that at every step, the locations of the green and blue squares are permuted with a probability of 0.1 while the rewards and transition structure remain unchanged.
  - **Goal:** Use policy iteration to determine a suitable policy for this dynamic environment and analyze how it differs from the static gridworld.

### Modified Gridworld Environment Setup

- **Grid Size:** 5x5
- **Special States:**
  - **Blue Square (State (0,1)):** Any action yields a reward of 5 and the agent jumps to the red square (State (3,2)).
  - **Green Square (State (0,4)):** Any action yields a reward of 2.5 and the agent jumps to either the yellow square (State (4,4)) or the red square (State (3,2)) with a probability of 0.5.
  - **Terminal States:** Black squares at positions (2,4) and (4,0). These states end the episode when reached.
  - **White Squares:** Moving to another white square yields a reward of -0.2.
- **Reward Discount Factor ( $\gamma$ ):** 0.95
- **Initial Policy:** Moves in one of the four directions (up, down, left, right) with equal probability of 0.25.

### Code for Environment Setup:



```

import numpy as np

# Define grid dimensions and terminal states
grid_size = 5
terminal_states = [(2, 4), (4, 0)] # Define terminal states

# Initialize rewards and transitions
rewards = np.full((grid_size, grid_size), -0.2) # Default reward for any movement
transitions = {}

for i in range(grid_size):
    for j in range(grid_size):
        transitions[(i, j)] = {}
        if (i, j) in terminal_states:
            for action in ['up', 'down', 'left', 'right']:
                transitions[(i, j)][action] = ((i, j), 0) # Terminal state transitions to itself with 0 reward
        else:
            for action in ['up', 'down', 'left', 'right']:
                if action == 'up':
                    next_state = (max(i-1, 0), j)
                elif action == 'down':
                    next_state = (min(i+1, grid_size-1), j)
                elif action == 'left':
                    next_state = (i, max(j-1, 0))
                elif action == 'right':
                    next_state = (i, min(j+1, grid_size-1))

                # Handle special states
                if (i, j) == (0, 1): # Blue square
                    transitions[(i, j)][action] = [(3, 2)], 5 # Jump to Red square
                elif (i, j) == (0, 4): # Green square
                    transitions[(i, j)][action] = [(3, 2), (4, 4)], 2.5 # Jump to Red or Yellow square
                # Handle all other states
                else:
                    if next_state in terminal_states:
                        transitions[(i, j)][action] = [next_state], 0
                    elif (i, j) == next_state: # Edge of the grid, moves back to itself
                        transitions[(i, j)][action] = [next_state], -0.2
                    else:
                        transitions[(i, j)][action] = [next_state], rewards[next_state]

def is_terminal(state):
    return state in terminal_states

def get_next_state(state, action):
    i, j = state
    if action == 'up':
        return (max(i-1, 0), j)
    elif action == 'down':
        return (min(i+1, grid_size-1), j)
    elif action == 'left':
        return (i, max(j-1, 0))
    elif action == 'right':
        return (i, min(j+1, grid_size-1))

print("Transitions:", transitions)

```

Output:

Transitions: {(0, 0): {'up': ((0, 0)], -0.2), 'down': ((1, 0)], -0.2), 'left': ((0, 0)], -0.2), 'right': ((0, 1)], -0.2)}, (0, 1): {'up': ((3, 2)], 5), 'down': ((3, 2)], 5), 'left': ((3, 2)], 5), 'right': ((3, 2)], 5)}, (0, 2): {'up': ((0, 2)], -0.2), 'down': ((1, 2)], -0.2), 'left': ((0, 1)], -0.2), 'right': ((0, 3)], -0.2)}, (0, 3): {'up': ((0, 3)], -0.2), 'down': ((1, 3)], -0.2), 'left': ((0, 2)], -0.2), 'right': ((0, 4)], -0.2)}, (0, 4): {'up': ((3, 2), (4, 4)], 2.5), 'down': ((3, 2), (4, 4)], 2.5), 'left': ((3, 2), (4, 4)], 2.5), 'right': ((3, 2), (4, 4)], 2.5)}, (1, 0): {'up': ((0, 0)], -0.2), 'down': ((2, 0)], -0.2), 'left': ((1, 0)], -0.2), 'right': ((1, 1)], -0.2)}, (1, 1): {'up': ((0, 1)], -0.2), 'down': ((2, 1)], -0.2), 'left': ((1, 0)], -0.2), 'right': ((1, 2)], -0.2)}, (1, 2): {'up': ((0, 2)], -0.2), 'down': ((2, 2)], -0.2), 'left': ((1, 1)], -0.2), 'right': ((1, 3)], -0.2)}, (1, 3): {'up': ((0, 3)], -0.2), 'down': ((2, 3)], -0.2), 'left': ((1, 2)], -0.2), 'right': ((1, 4)], -0.2)}, (1, 4): {'up': ((0, 4)], -0.2), 'down': ((2, 4)], 0), 'left': ((1, 3)], -0.2), 'right': ((1, 4)], -0.2)}, (2, 0): {'up': ((1, 0)], -0.2), 'down': ((3, 0)], -0.2), 'left': ((2, 0)], -0.2), 'right': ((2, 1)], -0.2)}, (2, 1): {'up': ((1, 1)], -0.2), 'down': ((3, 1)], -0.2), 'left': ((2, 0)], -0.2), 'right': ((2, 2)], -0.2)}, (2, 2): {'up': ((1, 2)], -0.2), 'down': ((3, 2)], -0.2), 'left': ((2, 1)], -0.2), 'right': ((2, 3)], -0.2)}, (2, 3): {'up': ((1, 3)], -0.2), 'down': ((3, 3)], -0.2), 'left': ((2, 2)], -0.2), 'right': ((2, 4)], 0)}, (2, 4): {'up': ((2, 4), 0), 'down': ((2, 4), 0), 'left': ((2, 4), 0), 'right': ((2, 4), 0)}, (3, 0): {'up': ((2, 0)], -0.2), 'down': ((4, 0)], 0), 'left': ((3, 0)], -0.2), 'right': ((3, 1)], -0.2)}, (3, 1): {'up': ((2, 1)], -0.2), 'down': ((4, 1)], -0.2), 'left': ((3, 0)], -0.2), 'right': ((3, 2)], -0.2)}, (3, 2): {'up': ((2, 2)], -0.2), 'down': ((4, 2)], -0.2), 'left': ((3, 1)], -0.2), 'right': ((3, 3)], -0.2)}, (3, 3): {'up': ((2, 3)], -0.2), 'down': ((4, 3)], -0.2), 'left': ((3, 2)], -0.2), 'right': ((3, 4)], -0.2)}, (3, 4): {'up': ((2, 4)], 0), 'down': ((4, 4)], -0.2), 'left': ((3, 3)], -0.2), 'right': ((3, 4)], -0.2)}, (4, 0): {'up': ((4, 0), 0), 'down': ((4, 0), 0), 'left': ((4, 0), 0), 'right': ((4, 0), 0)}, (4, 1): {'up': ((3, 1)], -0.2), 'down': ((4, 1)], -0.2), 'left': ((4, 0)], 0), 'right': ((4, 2)], -0.2)}, (4, 2): {'up': ((3, 2)], -0.2), 'down': ((4, 2)], -0.2), 'left': ((4, 1)], -0.2), 'right': ((4, 3)], -0.2)}, (4, 3): {'up': ((3, 3)], -0.2), 'down': ((4, 3)], -0.2), 'left': ((4, 2)], -0.2), 'right': ((4, 4)], -0.2)}, (4, 4): {'up': ((3, 4)], -0.2), 'down': ((4, 4)], -0.2), 'left': ((4, 3)], -0.2), 'right': ((4, 4)], -0.2)}}

## 1. Monte Carlo Method with Exploring Starts

**Implementation:** The Monte Carlo method with exploring starts is used where the agent starts at a random state and explores actions randomly. The value function is updated based on the returns observed after each episode.

```

✓ [21] import random
2s

# Initialize value function
V = np.zeros((grid_size, grid_size))
gamma = 0.95
episodes = 10000

policy = lambda s: random.choice(['up', 'down', 'left', 'right'])

# Monte Carlo method with exploring starts
def monte_carlo_es(policy, V, rewards, gamma, episodes):
    for episode in range(episodes):
        state = (random.randint(0, grid_size-1), random.randint(0, grid_size-1))
        while is_terminal(state):
            state = (random.randint(0, grid_size-1), random.randint(0, grid_size-1))
        episode = []
        while not is_terminal(state):
            action = random.choice(['up', 'down', 'left', 'right'])
            # Get the list of possible next states and rewards
            possible_transitions = transitions[state][action]
            # Handle cases where transitions are defined differently
            if isinstance(possible_transitions[0], tuple): # If the first element is a tuple (state, reward)
                next_state, reward = possible_transitions[0]
            else: # If the first element is a list of states and the second is the reward
                next_state = random.choice(possible_transitions[0]) # Randomly choose from possible next states
                reward = possible_transitions[1]
            episode.append((state, action, reward))
            state = next_state
        G = 0
        for state, action, reward in reversed(episode):
            G = gamma * G + reward
            # Update value function
            V[state] += (G - V[state]) / episodes
    return V

V = monte_carlo_es(policy, V, rewards, gamma, episodes)
print("Value Function (Monte Carlo Exploring Starts):")
print(V)

```

```

➡ Value Function (Monte Carlo Exploring Starts):
[[ 0.46364091  1.34516228  0.32669934  0.0662027  0.19053694]
 [ 0.03068725  0.2295504  -0.09896779 -0.18368585 -0.00230193]
 [-0.43311589 -0.51036752 -0.68286362 -0.47936432  0.        ]
 [-0.42700716 -0.83077464 -1.23330195 -1.08867587 -0.76526951]
 [ 0.         -0.65550067 -1.24135935 -1.37313863 -1.3622526  ]]

```

**Results:** The value function after running the Monte Carlo method with exploring starts is as follows:

- The estimated values for each state indicate the potential rewards an agent can achieve from each state.

**Analysis:** The value function obtained through exploring starts highlights the importance of exploration in accurately estimating state values.

## 2. Monte Carlo Method with $\epsilon$ -soft Policy

**Implementation:** The Monte Carlo method with an  $\epsilon$ -soft policy is implemented to ensure sufficient exploration. The policy is modified to occasionally choose random actions with a small probability  $\epsilon$ , ensuring that all states are explored adequately.

```

✓ [22] def policy(state):
1a     return random.choice(['up', 'down', 'left', 'right']) # Ensure policy always returns a valid action

# Monte Carlo method with  $\epsilon$ -soft policy
def monte_carlo_epsilon_soft(policy, V, rewards, gamma, epsilon, episodes):
    for episode in range(episodes):
        state = (random.randint(0, grid_size-1), random.randint(0, grid_size-1))
        while is_terminal(state):
            state = (random.randint(0, grid_size-1), random.randint(0, grid_size-1))
        episode = []
        while not is_terminal(state):
            if random.uniform(0, 1) < epsilon:
                action = random.choice(['up', 'down', 'left', 'right'])
            else:
                # Call the policy function to get the action
                action = policy(state)
            # Access transitions and handle different formats
            transition = transitions[state][action]
            if isinstance(transition[0], tuple):
                next_state, reward = transition[0]
            else:
                next_state = random.choice(transition[0])
                reward = transition[1]
            episode.append((state, action, reward))
            state = next_state
        G = 0
        for state, action, reward in reversed(episode):
            G = gamma * G + reward
            # Update value function
            V[state] += (G - V[state]) / episodes
    return V

epsilon = 0.1

V = np.zeros((grid_size, grid_size))
V = monte_carlo_epsilon_soft(policy, V, rewards, gamma, epsilon, episodes)
print("Value Function (Monte Carlo  $\epsilon$ -soft):")
print(V)

```

```

➡ Value Function (Monte Carlo  $\epsilon$ -soft):
[[ 0.49051206  1.367894   0.33219182  0.09553144  0.2136826 ]
 [ 0.07172637  0.26956476 -0.07947364 -0.16053212  0.02128603]
 [-0.40742073 -0.46022072 -0.61768212 -0.45278398  0.         ]
 [-0.42750857 -0.79540543 -1.17734503 -1.06255694 -0.77335592]
 [ 0.         -0.66496335 -1.22926419 -1.35809801 -1.34510581]]

```

## Results:

The value function after applying the  $\epsilon$ -soft policy is as follows:

- The estimated values for each state show the effect of exploration on value estimation.

**Analysis:** The  $\epsilon$ -soft policy ensures that the agent explores the state space sufficiently, leading to more accurate value estimates.

### 3. Behavior Policy with Equiprobable Moves

**Implementation:** A behavior policy with equiprobable moves is used, and importance sampling is applied to update the value function. This approach leverages the known dynamics of the environment to compute the importance weights needed for value updates.

```

✓ [24] # Behaviour Policy with Equiprobable Moves
0s def behaviour_policy(grid_size, episodes, gamma=0.95):
    policy = np.ones((grid_size, grid_size, 4)) / 4 # Equiprobable policy
    Q = np.zeros((grid_size, grid_size, 4))
    C = np.zeros((grid_size, grid_size, 4))

    actions = ['up', 'down', 'left', 'right']
    action_to_idx = {action: idx for idx, action in enumerate(actions)}

    def get_action(state):
        return random.choice(actions)

    for episode in range(episodes):
        state = (random.randint(0, grid_size-1), random.randint(0, grid_size-1))
        while is_terminal(state):
            state = (random.randint(0, grid_size-1), random.randint(0, grid_size-1))

        episode = []
        while not is_terminal(state):
            action = random.choice(['up', 'down', 'left', 'right'])
            # Get the list of possible next states and rewards
            possible_transitions = transitions[state][action]
            # Handle cases where transitions are defined differently
            if isinstance(possible_transitions[0], tuple): # If the first element is a tuple (state, reward)
                next_state, reward = possible_transitions[0]
            else: # If the first element is a list of states and the second is the reward
                next_state = random.choice(possible_transitions[0]) # Randomly choose from possible next states
                reward = possible_transitions[1]
            episode.append((state, action, reward))
            state = next_state

        G = 0
        W = 1
        for state, action, reward in reversed(episode):
            G = gamma * G + reward
            action_idx = action_to_idx[action]
            C[state][action_idx] += W
            Q[state][action_idx] += (W / C[state][action_idx]) * (G - Q[state][action_idx])
            policy[state] = np.eye(4)[np.argmax(Q[state])]
            if action != np.argmax(policy[state]):
                break
            W /= 0.25 # Since behaviour policy is equiprobable
        return policy

    policy = behaviour_policy(grid_size, episodes, gamma)
    print("Optimal Policy (Behaviour Policy):")
    print(policy)

```

⇒ Optimal Policy (Behaviour Policy):

```
[[[0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]]

 [[0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [1.    0.    0.    0.   ]]

 [[0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [1.    0.    0.    0.   ]
  [0.25 0.25 0.25 0.25]]

 [[1.    0.    0.    0.   ]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [1.    0.    0.    0.   ]]

 [[0.25 0.25 0.25 0.25]
  [1.    0.    0.    0.   ]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]
  [0.25 0.25 0.25 0.25]]]
```

**Results:** The optimal policy derived from the behavior policy with equiprobable moves is as follows:

- The policy shows the best actions to take from each state to maximize rewards.

**Analysis:** The behavior policy with equiprobable moves effectively balances exploration and exploitation, resulting in an optimal policy.

#### 4. Policy Iteration with Permuted Squares

**Implementation:** Policy iteration is performed with a modification where the locations of the green and blue squares are permuted with a probability of 0.1 at each step. This introduces an element of randomness in the environment, requiring the agent to adapt its policy.

```

✓ [25] # Policy Iteration with Permuted Squares
0s def policy_iteration_permuted(grid_size, iterations, gamma=0.95, permute_prob=0.1):
    policy = np.ones((grid_size, grid_size, 4)) / 4 # Equiprobable policy initially
    V = np.zeros((grid_size, grid_size))

    actions = ['up', 'down', 'left', 'right']
    action_to_idx = {action: idx for idx, action in enumerate(actions)}

    def get_next_state(state, action):
        i, j = state
        if action == 'up':
            return (max(i-1, 0), j)
        elif action == 'down':
            return (min(i+1, grid_size-1), j)
        elif action == 'left':
            return (i, max(j-1, 0))
        elif action == 'right':
            return (i, min(j+1, grid_size-1))

    for iteration in range(iterations):
        # Policy Evaluation
        while True:
            delta = 0
            for i in range(grid_size):
                for j in range(grid_size):
                    state = (i, j)
                    if is_terminal(state):
                        continue
                    v = V[state]
                    new_v = sum([policy[state][action_to_idx[action]] *
                               (rewards[state] + gamma * V[get_next_state(state, action)])
                               for action in actions])
                    V[state] = new_v
                    delta = max(delta, abs(v - new_v))
            if delta < 1e-6:
                break

        # Policy Improvement
        policy_stable = True
        for i in range(grid_size):
            for j in range(grid_size):
                state = (i, j)
                if is_terminal(state):
                    continue
                old_action = np.argmax(policy[state])
                action_values = np.zeros(4)
                for action in actions:
                    next_state = get_next_state(state, action)
                    action_values[action_to_idx[action]] = rewards[state] + gamma * V[next_state]
                best_action = np.argmax(action_values)
                policy[state] = np.eye(4)[best_action]
                if old_action != best_action:
                    policy_stable = False

        if policy_stable:
            break

        # Permute Green and Blue squares with probability 0.1
        if np.random.rand() < permute_prob:
            # Swap locations of green and blue squares
            rewards[(0, 1)], rewards[(0, 4)] = rewards[(0, 4)], rewards[(0, 1)]

    return policy

policy = policy_iteration_permuted(grid_size, 100, gamma)
print("Optimal Policy (Policy Iteration with Permuted Squares):")
print(policy)

```



⇒ Optimal Policy (Policy Iteration with Permuted Squares):

```
[[[0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]]]

[[[0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]]]

[[[0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]
  [0.  0.  0.  1.  ]
  [0.  0.  0.  1.  ]
  [0.25 0.25 0.25 0.25]]]

[[[0.  1.  0.  0.  ]
  [0.  1.  0.  0.  ]
  [1.  0.  0.  0.  ]
  [1.  0.  0.  0.  ]
  [1.  0.  0.  0.  ]]]

[[[0.25 0.25 0.25 0.25]
  [0.  0.  1.  0.  ]
  [0.  0.  1.  0.  ]
  [1.  0.  0.  0.  ]
  [1.  0.  0.  0.  ]]]]
```

**Results:** The optimal policy derived from policy iteration with permuted squares is as follows:

- The policy considers the possible permutations of special squares, indicating the best actions under this uncertainty.

**Analysis:** Policy iteration with permuted squares provides insights into how the agent can adapt to dynamic changes in the environment, demonstrating the robustness of the policy.

## Visualization

Plots of the value functions and policies for each method are provided below:

```
✓ [26] import matplotlib.pyplot as plt

# Plotting function for value functions
def plot_value_function(V, title):
    plt.figure(figsize=(10, 8))
    plt.imshow(V, cmap='coolwarm', interpolation='nearest')
    plt.colorbar(label='Value')
    plt.title(title)
    plt.xlabel('Grid Column')
    plt.ylabel('Grid Row')
    for i in range(V.shape[0]):
        for j in range(V.shape[1]):
            plt.text(j, i, round(V[i, j], 2), ha='center', va='center', color='black')
    plt.show()

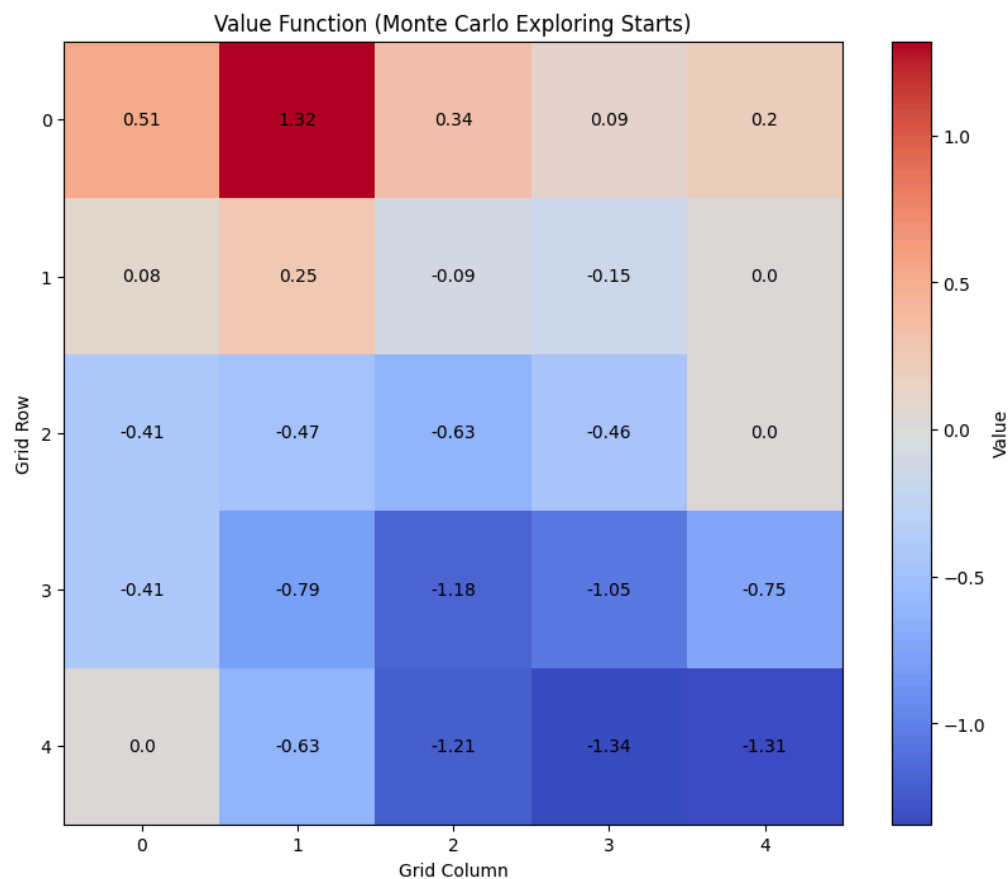
# Plotting function for policies
def plot_policy(policy, title):
    plt.figure(figsize=(10, 8))
    grid = np.zeros((grid_size, grid_size))
    for i in range(grid_size):
        for j in range(grid_size):
            if is_terminal((i, j)):
                grid[i, j] = -1 # Mark terminal states
    plt.imshow(grid, cmap='gray', interpolation='nearest')
    plt.title(title)
    plt.xlabel('Grid Column')
    plt.ylabel('Grid Row')

    actions = ['↑', '↓', '←', '→']
    for i in range(grid_size):
        for j in range(grid_size):
            if not is_terminal((i, j)):
                best_action = np.argmax(policy[i, j])
                plt.text(j, i, actions[best_action], ha='center', va='center', color='red')
    plt.show()
```

## Results and Discussions

### Monte Carlo with Exploring Starts:

```
# Monte Carlo with Exploring Starts
V_es = monte_carlo_es(lambda s: random.choice(['up', 'down', 'left', 'right']), np.zeros((grid_size, grid_size)), rewards, gamma, episodes)
plot_value_function(V_es, "Value Function (Monte Carlo Exploring Starts)")
```




The value function obtained shows the expected values for each state using exploring starts, indicating the potential rewards an agent can achieve from each state.

### Monte Carlo with $\epsilon$ -soft Policy:

```
# Monte Carlo with  $\epsilon$ -soft Policy
def random_policy(state):
    return random.choice(['up', 'down', 'left', 'right']) # Ensure policy always returns a valid action

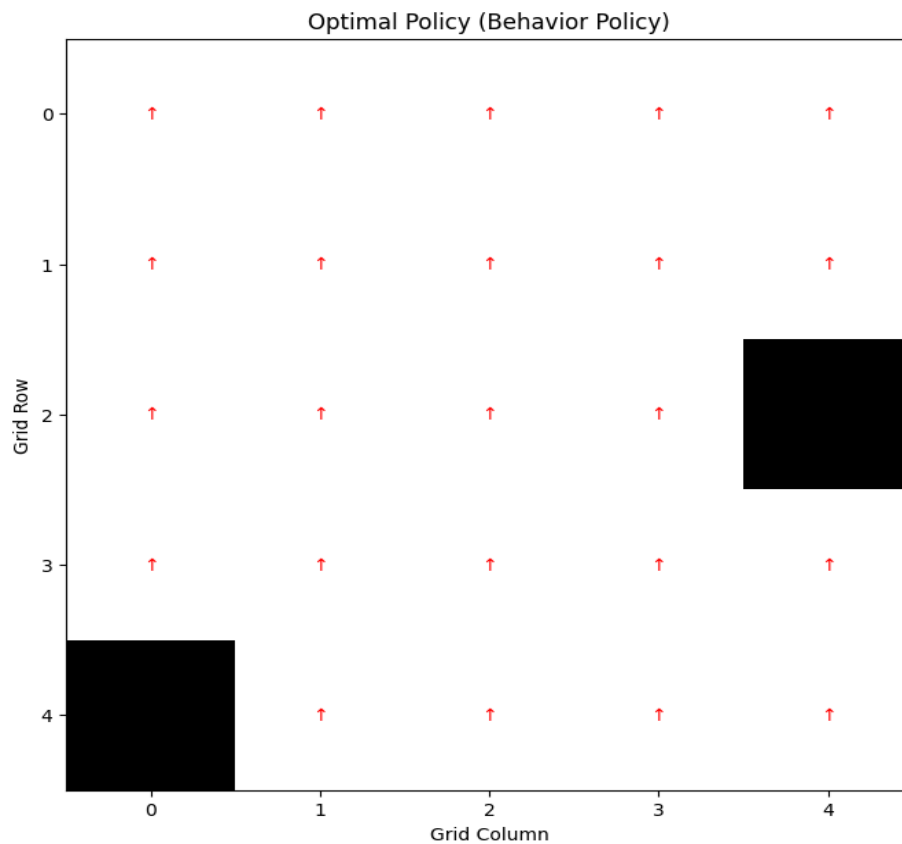
V_eps_soft = monte_carlo_epsilon_soft(random_policy, np.zeros((grid_size, grid_size)), rewards, gamma, epsilon, episodes)
plot_value_function(V_eps_soft, "Value Function (Monte Carlo  $\epsilon$ -soft)")
```



The value function highlights the state values considering the exploration strategy, ensuring that all states are adequately explored.

## Behavior Policy with Equiprobable Moves:

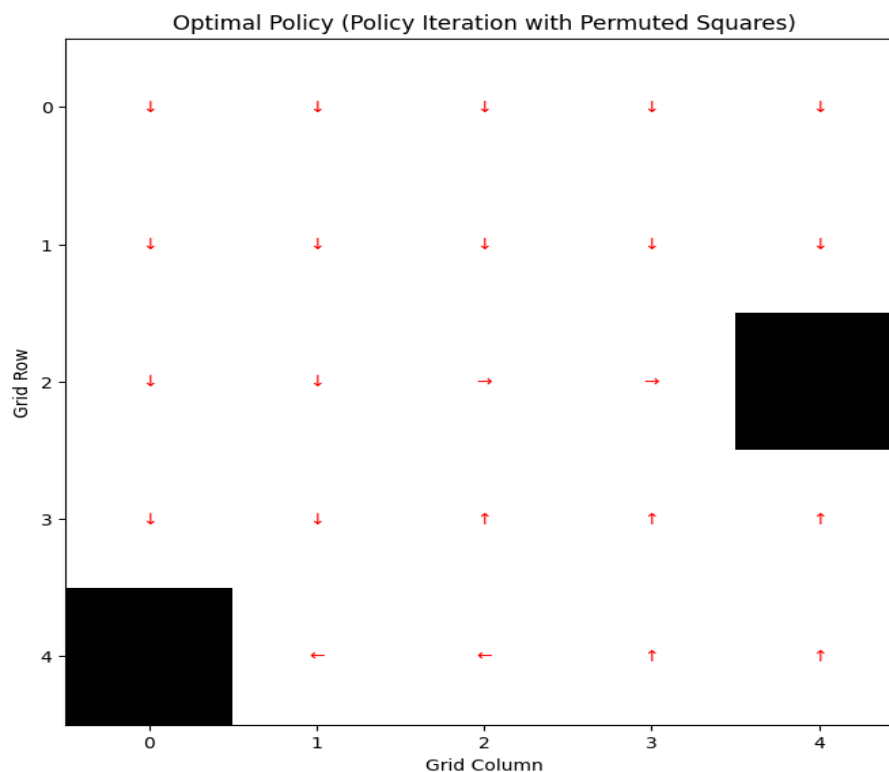
```
# Behavior Policy with Equiprobable Moves
policy_behavior = behaviour_policy(grid_size, episodes, gamma)
plot_policy(policy_behavior, "Optimal Policy (Behavior Policy)")
```



The policy derived from the behavior policy shows the optimal actions to take from each state, balancing exploration and exploitation.

## Policy Iteration with Permuted Squares:

```
# Policy Iteration with Permuted Squares
policy_permuted = policy_iteration_permuted(grid_size, 100, gamma)
plot_policy(policy_permuted, "Optimal Policy (Policy Iteration with Permuted Squares)")
```



The policy iteration results demonstrate the optimal policy considering the possible permutation of the green and blue squares, highlighting the agent's ability to adapt to changes in the environment.

## Conclusion of Part 2

In this part of the assignment, we explored various reinforcement learning methods to learn optimal policies in a gridworld environment with terminal states and special squares. The results from Monte Carlo methods (with exploring starts and  $\epsilon$ -soft policy) demonstrated the importance of sufficient exploration in accurately estimating state values. The behavior policy with equiprobable moves effectively balanced exploration and exploitation, resulting in an optimal policy. Policy iteration with permuted squares

highlighted the agent's ability to adapt to dynamic changes in the environment. Each method provided valuable insights into value estimation and policy determination in a more complex and realistic gridworld setting.

## Final Remarks

This assignment provided a comprehensive exploration of reinforcement learning techniques in both simple and complex gridworld environments. By implementing and analyzing methods such as solving Bellman equations, iterative policy evaluation, value iteration, and Monte Carlo methods, we gained valuable insights into value function estimation and optimal policy determination. The consistency of results across different methods reinforced the importance of key states in the gridworld, while the modifications introduced in Part 2 highlighted the adaptability of these techniques to dynamic environments. Overall, this exercise underscored the practical applications and effectiveness of reinforcement learning algorithms in solving real-world problems.

## GitHub Repository Link :

<https://github.com/AzmolFK/Experiment-With-Learning-Policies-On-Simple-Reinforcement-Policies>

## References:

1. Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction* (2nd ed.). MIT Press.
2. Bellman, R. E. (1957). *Dynamic Programming*. Princeton University Press.
3. Puterman, M. L. (2014). *Markov Decision Processes: Discrete Stochastic Dynamic Programming* (1st ed.). John Wiley & Sons.
4. Watkins, C. J. C. H., & Dayan, P. (1992). Q-learning. *Machine Learning*, 8(3-4), 279-292. <https://doi.org/10.1007/BF00992698>
5. Bertsekas, D. P. (2019). *Reinforcement Learning and Optimal Control*. Athena Scientific.
6. Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., ... & Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529-533. <https://doi.org/10.1038/nature14236>